

Stage 3: Database Implementation and Indexing (22%)

Implement at least four main tables

- User
- Health_record
- Food
- Exercise
- Order_records

```
mysql> show databases;
+-----+
| Database |
+-----+
| information_schema |
| classicmodels |
| mysql |
| performance_schema |
| sys |
+-----+
5 rows in set (0.01 sec)
```

```
mysql> use classicmodels;
Database changed
```

```
Database changed
mysql> show tables;
+-----+
| Tables_in_classicmodels |
+-----+
| Exercise                 |
| Food                     |
| health_record            |
| order_record             |
| user                     |
+-----+
5 rows in set (0.00 sec)
```

In the Database Design markdown or pdf, provide the Data Definition Language (DDL) commands you all used to create each of these tables in the database.

Markdown file:

<https://s3-us-west-2.amazonaws.com/secure.notion-static.com/2bbe5576-e7a5-4b8c-9825-34b0754a2e6d/ddl.md>

Insert data to these tables. You should insert at least 1000 rows each in three of the tables.

- User
- Health_record
- Food

```
mysql> SELECT COUNT(foodID) as countFood
-> FROM Food
-> ORDER BY foodID;
+-----+
| countFood |
+-----+
|      7083 |
+-----+
1 row in set (0.00 sec)
```

```
mysql> SELECT COUNT(healthID) as countHealth FROM health_record ORDER BY healthID;
+-----+
| countHealth |
+-----+
|      1000 |
+-----+
1 row in set (0.00 sec)
```

```
mysql> SELECT COUNT(userID) as countUser FROM user ORDER BY userID;
+-----+
| countUser |
+-----+
|      1000 |
+-----+
1 row in set (0.06 sec)
```

Datas inside our tables

```
mysql> SELECT * FROM user LIMIT 15;
```

userID	firstName	lastName	email	pass
1	Dannie	Garnham	dgarnham0@comcast.net	6o251
2	Caryl	Sleight	csleight1@i2i.jp	1BR54
3	Vaughn	Brigginshaw	vbrigginshaw2@squarespace.com	1Ef59
4	Carlin	Hemeret	chemeret3@youku.com	LWi49
5	Ode	Huston	ohuston4@soundcloud.com	QZL35
6	Sunshine	Philippon	sphilippon5@lycos.com	U6m62
7	Hyacinthie	Brantl	hbrantl6@nsw.gov.au	gqH25
8	Curry	Roose	croose7@so-net.ne.jp	AUh75
9	Fidole	Kahan	fkahan8@prnewswire.com	Lua52
10	Starlin	Pfleger	spfleger9@booking.com	RSn94
11	Romy	Taggert	rtaggerta@webs.com	fx165
12	Lola	Sommerlie	lsommerlieb@ebay.com	cIi61
13	Ephraim	Chalke	echalkec@thetimes.co.uk	lp105
14	Tess	Kingsly	tkingslyd@jiathis.com	Rfy35
15	Helga	Rubes	hrubese@psu.edu	qK922

```
15 rows in set (0.01 sec)
```

```
mysql> SELECT * FROM health_record LIMIT 15;
```

healthID	healthUserID	gender	weight	height	BMI	BMR	CaloriesNeeded	curr_date
710784	550	M	97	60	7	15	7147	12/01/2022
20076398	271	M	145	83	21	20	9947	02/09/2021
24388491	734	M	77	149	61	15	9811	25/01/2022
26782553	505	M	86	133	71	92	2263	16/09/2021
26957221	406	F	107	54	76	99	2796	28/01/2022
49096818	294	M	88	131	76	82	1619	27/07/2021
50853414	260	M	66	129	89	74	3976	09/02/2022
53065034	92	F	145	95	79	85	9880	01/02/2022
61510971	50	F	134	95	67	43	8700	31/08/2021
63203960	336	F	84	137	70	25	6783	12/06/2021
78934702	350	M	85	124	24	17	4857	22/11/2021
79659802	956	M	133	97	19	70	382	24/04/2021
88485749	912	F	98	109	45	73	8359	02/01/2022
103536795	227	M	83	121	58	39	5008	20/12/2021
123861241	738	M	112	64	47	9	9188	19/08/2021

```
15 rows in set (0.00 sec)
```

```
mysql> SELECT * FROM Exercise LIMIT 15;
```

exerciseID	exerciseName	exerciseType
1000	Aerobics - Low impact	Low intensity
1001	Aerobics - Step	High intensity
1002	Archery	Cardio
1003	Badminton	Cardio
1004	'Ballet, twist, jazz, tap'	Cardio
1005	Basketball - Shooting Hoops	High intensity
1006	Bicycling - 12 - 14 mph moderate	High intensity
1007	Bike - Stationary	High intensity
1008	Bowling	Low intensity
1009	Boxing - Punching bag	High intensity
1010	'Carrying 16 to 24 lbs, upstairs'	High intensity
1011	Chopping Wood	High intensity
1012	Circuit Training	High intensity
1013	Cleaning House	Low intensity
1014	Climbing hills	High intensity

```
15 rows in set (0.00 sec)
```

```
mysql> SELECT * FROM Food LIMIT 15;
```

foodID	foodName	carbsCalories	proteinCalories	fiberCalories
11000000	Milk, whole	70.00	6.89	1.03
11100000	Milk, low sodium, whole	51.00	4.87	3.34
11111000	Milk, calcium fortified, whole	60.00	4.67	3.28
11111100	Milk, calcium fortified, low fat (1%)	61.00	4.46	3.10
11111150	Milk, calcium fortified, fat free (skim)	60.00	4.67	3.28
11111160	Milk, reduced fat (2%)	43.00	5.19	3.38
11111170	Milk, acidophilus, low fat (1%)	35.00	4.85	3.40
11112110	Milk, acidophilus, reduced fat (2%)	50.00	4.91	3.35
11112120	Milk, low fat (1%)	43.00	5.19	3.38
11112130	Milk, fat free (skim)	50.00	4.91	3.35
11112210	Milk, lactose free, low fat (1%)	43.00	5.19	3.38
11113000	Milk, lactose free, fat free (skim)	34.00	4.89	3.43
11114300	Milk, lactose free, reduced fat (2%)	43.00	5.19	3.38
11114320	Milk, lactose free, whole	34.00	4.89	3.43
11114330	Buttermilk, fat free (skim)	50.00	4.91	3.35

```
15 rows in set (0.00 sec)
```

order_record is empty and datas will be added once the program runs.

2 advanced SQL Queries:

1st Advanced SQL Query:

→ The Categories filter

In our website there will be a filter in a form of checkbox where it can display foods that has high carbs/protein/fiber content. It will be something like:

Categories

- ☐ High Carbohydrates
- ☐ High Protein
- ☐ High Fiber

Our first Advanced SQL Query will be when an instance where the user pick 2 filters, like:

Categories

- ☒ High Carbohydrates
- ☒ High Protein
- ☐ High Fiber

```
SELECT f1.foodName, f1.carbsCalories, f1.proteinCalories, f1.fiberCalories
FROM Food f1
WHERE f1.carbsCalories > (SELECT AVG(f2.carbsCalories)
                        FROM Food f2))
AND f1.foodName IN
(SELECT f3.foodName
FROM Food f3
WHERE f3.proteinCalories > (SELECT AVG(f4.proteinCalories)
                        FROM Food f4));
```

This is an advanced SQL Query because it uses subquery and set operation (INTERSECT), which is replaced by AND ____ IN since to achieve the INTERSECT operation functionality since INTERSECT is not available in MYSQL.

Below are another cases which is similar to the first example:

Categories

- ☐ High Carbohydrates
- ☒ High Protein
- ☒ High Fiber

```
(SELECT f1.foodName, f1.carbsCalories, f1.proteinCalories, f1.fiberCalories
FROM Food f1
WHERE f1.proteinCalories > (SELECT AVG(f2.proteinCalories)
                           FROM Food f2))
AND f1.foodName IN
(SELECT f3.foodName, f3.carbsCalories, f3.proteinCalories, f3.fiberCalories
FROM Food f3
WHERE f3.fiberCalories > (SELECT AVG(f4.fiberCalories)
                           FROM Food f4))
ORDER BY foodName;
```

Categories

☒ High Carbohydrates

☐ High Protein

☒ High Fiber

```
(SELECT f1.foodName, f1.carbsCalories, f1.proteinCalories, f1.fiberCalories
FROM Food f1
WHERE f1.carbsCalories > (SELECT AVG(f2.carbsCalories)
                           FROM Food f2))
WHERE f1.foodName, f1.carbsCalories, f1.proteinCalories, f1.fiberCalories IN
(SELECT f3.foodName, f3.carbsCalories, f3.proteinCalories, f3.fiberCalories
FROM Food f3
WHERE f3.fiberCalories > (SELECT AVG(f4.fiberCalories)
                           FROM Food f4))
ORDER BY foodName;
```

Categories

☒ Carbohydrates

☒ Protein

☒ Fiber

```
(SELECT f1.foodName, f1.carbsCalories, f1.proteinCalories, f1.fiberCalories
FROM Food f1
WHERE f1.carbsCalories > (SELECT AVG(f2.carbsCalories)
                           FROM Food f2))
AND f1.foodName IN
(SELECT f3.foodName, f3.carbsCalories, f3.proteinCalories, f3.fiberCalories
FROM Food f3
WHERE f3.fiberCalories > (SELECT AVG(f4.fiberCalories)
```

```

FROM Food f4))

INTERSECT
(SELECT f5.foodName, f5.carbsCalories, f5.proteinCalories, f5.fiberCalories
FROM Food f5
WHERE f6.proteinCalories > (SELECT AVG(f6.proteinCalories)
FROM Food f6))

ORDER BY foodName;

```

2nd Advanced SQL Query:

→ This query returns the history of a user's BMI. We are querying healthUserID = 1 as an example.

```

SELECT firstName, lastName, MAX(BMI) as maxBMI, MIN(BMI) as minBMI, AVG(BMI) as avgBMI
FROM health_record JOIN user ON (healthUserID = userID)
GROUP BY healthUserID
HAVING healthUserID = 1

```

We are joining based on userID, which is shared amongst health_record and user tables since we want to display the first and last name of the user and their BMI history on the screen. First and last name of the user information are retrieved from user table, whereas the BMI information is retrieved from health_record table.

Since health_record is auto-generated, we insert some datas into the health_record database to mimic a returning user. We are inserting to healthUserID = 1 as an example.

```

INSERT INTO health_record VALUES
(0,1,'M',146,76,18,5,2679,'16/02/2022'),
(1,1,'M',146,76,18.1,5,2679,'17/02/2022'),
(2,1,'M',146,76,18.2,5,2679,'18/02/2022'),
(3,1,'M',146,76,18.4,5,2679,'19/02/2022'),
(4,1,'M',146,76,18.5,5,2679,'20/02/2022'),
(5,1,'M',146,76,18.8,5,2679,'21/02/2022'),
(6,1,'M',146,76,18.9,5,2679,'22/02/2022'),
(7,1,'M',146,76,18.3,5,2679,'23/02/2022'),
(8,1,'M',146,76,18.4,5,2679,'24/02/2022'),
(9,1,'M',146,76,18.8,5,2679,'25/02/2022'),
(10,1,'M',146,76,18.4,5,2679,'26/02/2022'),
(11,1,'M',146,76,18,5,2679,'27/02/2022'),
(12,1,'M',146,76,18.2,5,2679,'28/02/2022'),

```



```
(13,1,'M',146,76,20,5,2679,'1/03/2022'),
(14,1,'M',146,76,21,5,2679,'2/03/2022'),
(15,1,'M',146,76,22,5,2679,'3/03/2022'),
(16,1,'M',146,76,23,5,2679,'4/03/2022'),
(17,1,'M',146,76,22,5,2679,'5/03/2022'),
(18,1,'M',146,76,21,5,2679,'6/03/2022'),
(19,1,'M',146,76,22,5,2679,'7/03/2022'),
(20,1,'M',146,76,23,5,2679,'8/03/2022'),
(21,1,'M',146,76,20,5,2679,'9/03/2022'),
(22,1,'M',146,76,23,5,2679,'10/03/2022'),
(23,1,'M',146,76,25,5,2679,'11/03/2022'),
(24,1,'M',146,76,22,5,2679,'12/03/2022'),
(25,1,'M',146,76,22,5,2679,'13/03/2022'),
(26,1,'M',146,76,20,5,2679,'14/03/2022'),
(27,1,'M',146,76,18.3,5,2679,'15/03/2022'),
(28,1,'M',146,76,18.5,5,2679,'16/03/2022'),
(29,1,'M',146,76,18.8,5,2679,'17/03/2022'),
(30,1,'M',146,76,19,5,2679,'18/03/2022'),
(31,1,'M',146,76,19.1,5,2679,'19/03/2022'),
(32,1,'M',146,76,18.3,5,2679,'20/03/2022'),
(33,1,'M',146,76,18.4,5,2679,'21/03/2022'),
(34,1,'M',146,76,19.6,5,2679,'22/03/2022'),
(35,1,'M',146,76,19.9,5,2679,'23/03/2022'),
(36,1,'M',146,76,19.8,5,2679,'24/03/2022'),
(37,1,'M',146,76,19.8,5,2679,'25/03/2022'),
(38,1,'M',146,76,18.8,5,2679,'26/03/2022'),
(39,1,'M',146,76,17.8,5,2679,'27/03/2022'),
(40,1,'M',146,76,17.8,5,2679,'28/03/2022'),
(41,1,'M',146,76,18.2,5,2679,'29/03/2022'),
(42,1,'M',146,76,18.9,5,2679,'30/03/2022'),
(43,1,'M',146,76,18.8,5,2679,'31/03/2022'),
(44,1,'M',146,76,18.5,5,2679,'1/04/2022'),
(45,1,'M',146,76,18.9,5,2679,'2/04/2022'),
(46,1,'M',146,76,18.4,5,2679,'3/04/2022'),
(47,1,'M',146,76,18.7,5,2679,'4/04/2022'),
(48,1,'M',146,76,16,5,2679,'5/04/2022'),
(49,1,'M',146,76,16,5,2679,'6/04/2022'),
(50,1,'M',146,76,16.5,5,2679,'7/04/2022');
```

Checking if the insertion is successful:

```
mysql> SELECT * FROM health_record WHERE healthUserID = 1;
```

healthID	healthUserID	gender	weight	height	BMI	BMR	CaloriesNeeded	curr_date
0	1	M	146	76	18	5	2679	16/02/2022
1	1	M	146	76	18	5	2679	17/02/2022
2	1	M	146	76	18	5	2679	18/02/2022
3	1	M	146	76	18	5	2679	19/02/2022
4	1	M	146	76	19	5	2679	20/02/2022
5	1	M	146	76	19	5	2679	21/02/2022
6	1	M	146	76	19	5	2679	22/02/2022
7	1	M	146	76	18	5	2679	23/02/2022
8	1	M	146	76	18	5	2679	24/02/2022
9	1	M	146	76	19	5	2679	25/02/2022
10	1	M	146	76	18	5	2679	26/02/2022
11	1	M	146	76	18	5	2679	27/02/2022
12	1	M	146	76	18	5	2679	28/02/2022
13	1	M	146	76	20	5	2679	1/03/2022
14	1	M	146	76	21	5	2679	2/03/2022
15	1	M	146	76	22	5	2679	3/03/2022
16	1	M	146	76	23	5	2679	4/03/2022
17	1	M	146	76	22	5	2679	5/03/2022
18	1	M	146	76	21	5	2679	6/03/2022
19	1	M	146	76	22	5	2679	7/03/2022
20	1	M	146	76	23	5	2679	8/03/2022
21	1	M	146	76	20	5	2679	9/03/2022
22	1	M	146	76	23	5	2679	10/03/2022
23	1	M	146	76	25	5	2679	11/03/2022
24	1	M	146	76	22	5	2679	12/03/2022
25	1	M	146	76	22	5	2679	13/03/2022
26	1	M	146	76	20	5	2679	14/03/2022
27	1	M	146	76	18	5	2679	15/03/2022
28	1	M	146	76	19	5	2679	16/03/2022

This is an advanced Query because we join relations and contains an aggregation via GROUP BY.

Execute your advanced SQL queries and provide a screenshot of the top 15 rows of each query result (you can use the LIMIT clause to select the top 15 rows).

— ss of top 15 of Advanced Query #1 —

```
mysql> SELECT f1.foodName, f1.carbsCalories, f1.proteinCalories, f1.fiberCalories FROM Food f1 WHERE f1.carbsCalories > (SELECT AVG(f2.carbsCalories) FROM Food f2) AND f1.foodName IN (SELECT f3.foodName FROM Food f3 WHERE f3.proteinCalories > (SELECT AVG(f4.proteinCalories) FROM Food f4)) limit 15;
```

foodName	carbsCalories	proteinCalories	fiberCalories
Soy milk	321.00	54.40	7.91
Frozen yogurt bar, chocolate	204.00	34.38	3.90
Milk, dry, not reconstituted, low fat (1%)	362.00	51.98	36.16
Milk, dry, not reconstituted, fat free (skim)	496.00	38.42	26.32
Whey, sweet, dry	367.00	51.46	35.80
Cocoa powder, not reconstituted	362.00	51.98	36.16
Chocolate beverage powder, dry mix, not reconstituted	353.00	74.46	12.93
Chocolate beverage powder, light, dry mix, not reconstituted	228.00	57.90	19.60
Milk, malted, dry mix, not reconstituted	400.00	90.28	4.55
Strawberry beverage powder, dry mix, not reconstituted	350.00	68.31	9.09
Cream, NS as to light, heavy, or half and half	372.00	92.96	0.00
Cream, light	389.00	99.10	0.10
Coffee creamer, liquid, fat free, flavored	251.00	35.07	0.69
Coffee creamer, powder, fat free	529.00	59.29	2.48
Coffee creamer, powder, fat free, flavored	482.00	75.42	0.68

15 rows in set (0.01 sec)

— ss of top 15 of Advanced Query #2—

```
mysql> SELECT firstName, lastName, MAX(BMI) as maxBMI, MIN(BMI) as minBMI, AVG(BMI) as avgBMI
-> FROM health_record JOIN user ON (healthUserID = userID)
-> GROUP BY healthUserID
-> HAVING healthUserID = 1;
+-----+-----+-----+-----+-----+
| firstName | lastName | maxBMI | minBMI | avgBMI |
+-----+-----+-----+-----+-----+
| Dannie   | Garnham  | 25     | 8      | 19.1731 |
+-----+-----+-----+-----+-----+
1 row in set (0.01 sec)
```

Advanced Query #2 only returns one row because it is querying a specific user's health records history to retrieve their maximum, minimum, and average BMI. The information we show on our website will be personalized according to which user is logged in, hence it makes sense to query one user at a time. In this case we are using healthUserID = 1 as an example.

INDEXING

Advanced Query One:

```
EXPLAIN ANALYZE
SELECT f1.foodName, f1.carbsCalories, f1.proteinCalories, f1.fiberCalories
FROM Food f1
WHERE f1.carbsCalories > (SELECT AVG(f2.carbsCalories)
                        FROM Food f2)
AND f1.foodName IN (SELECT f3.foodName
                   FROM Food f3
                   WHERE f3.proteinCalories > (SELECT AVG(f4.proteinCalories)
                                              FROM Food f4));
```

Index #1:

Without indexing:

```
+-----+
| -> Nested loop inner join (cost=586515.26 rows=5860069) (actual time=7.252..11.768 rows=1890 loops=1)
|   -> Filter: (f1.carbsCalories > (select #2)) (cost=266.33 rows=2421) (actual time=2.179..4.676 rows=2974 loops=1)
|     -> Table scan on f1 (cost=266.33 rows=7263) (actual time=0.043..1.697 rows=7083 loops=1)
|     -> Select #2 (subquery in condition; run only once)
|       -> Aggregate: avg(f2.carbsCalories) (cost=1476.85 rows=7263) (actual time=2.119..2.119 rows=1 loops=1)
|         -> Table scan on f2 (cost=750.55 rows=7263) (actual time=0.013..1.346 rows=7083 loops=1)
|   -> Single-row index lookup on <subquery3> using <auto distinct key> (foodName=f1.foodName) (actual time=0.000..0.000 rows=1 loops=2974)
|   -> Materialize with deduplication (cost=508.40..508.40 rows=2421) (actual time=6.478..6.657 rows=2270 loops=1)
|     -> Filter: (f3.proteinCalories > (select #4)) (cost=266.33 rows=2421) (actual time=1.935..4.226 rows=2270 loops=1)
|       -> Table scan on f3 (cost=266.33 rows=7263) (actual time=0.017..1.449 rows=7083 loops=1)
|       -> Select #4 (subquery in condition; run only once)
|         -> Aggregate: avg(f4.proteinCalories) (cost=1476.85 rows=7263) (actual time=1.907..1.907 rows=1 loops=1)
|           -> Table scan on f4 (cost=750.55 rows=7263) (actual time=0.015..1.232 rows=7083 loops=1)
|
+-----+
1 row in set (0.02 sec)
```

Creating an index on foodName:

```
mysql> CREATE INDEX foodName_idx on Food(foodName);
Query OK, 0 rows affected (0.11 sec)
Records: 0  Duplicates: 0  Warnings: 0
```

After indexing:

```
--+
| -> Nested loop semijoin (cost=1135.37 rows=828) (actual time=5.953..22.596 rows=1890 loops=1)
|   -> Filter: (f1.carbsCalories > (select #2)) (cost=266.33 rows=2421) (actual time=3.109..6.913 rows=2974 loops=1)
|     -> Table scan on f1 (cost=266.33 rows=7263) (actual time=0.045..2.202 rows=7083 loops=1)
|     -> Select #2 (subquery in condition; run only once)
|       -> Aggregate: avg(f2.carbsCalories) (cost=1476.85 rows=7263) (actual time=3.040..3.040 rows=1 loops=1)
|         -> Table scan on f2 (cost=750.55 rows=7263) (actual time=0.048..1.620 rows=7083 loops=1)
|   -> Filter: (f3.proteinCalories > (select #4)) (cost=0.09 rows=0) (actual time=0.005..0.005 rows=1 loops=2974)
|     -> Index lookup on f3 using foodName_idx (foodName=f1.foodName) (cost=0.09 rows=1) (actual time=0.004..0.004 rows=1 loops=2974)
|     -> Select #4 (subquery in condition; run only once)
|       -> Aggregate: avg(f4.proteinCalories) (cost=1476.85 rows=7263) (actual time=2.802..2.802 rows=1 loops=1)
|         -> Table scan on f4 (cost=750.55 rows=7263) (actual time=0.015..1.421 rows=7083 loops=1)
+-----+
```

Without the indexing, the cost of the nested loop inner join is 586515.26 and it has to get through 5860069 rows, after creating an index for the food name (foodName_idx) the cost significantly drops to 1135.37 and it only has to go through 828 rows.

```
--+
1 row in set (0.03 sec)

mysql> CREATE INDEX carbsCal_idx on Food(carbsCalories);
Query OK, 0 rows affected (0.07 sec)
Records: 0  Duplicates: 0  Warnings: 0
```

```
+-----+
| -> Nested loop semijoin (cost=1818.20 rows=1017) (actual time=2.897..16.668 rows=1890 loops=1)
|   -> Filter: (f1.carbsCalories > (select #2)) (cost=750.55 rows=2974) (actual time=0.041..4.086 rows=2974 loops=1)
|     -> Table scan on f1 (cost=750.55 rows=7263) (actual time=0.029..2.296 rows=7083 loops=1)
|     -> Select #2 (subquery in condition; run only once)
|       -> Aggregate: avg(f2.carbsCalories) (cost=1476.85 rows=7263) (actual time=2.790..2.790 rows=1 loops=1)
|         -> Index scan on f2 using carbsCal_idx (cost=750.55 rows=7263) (actual time=0.037..1.343 rows=7083 loops=1)
|   -> Filter: (f3.proteinCalories > (select #4)) (cost=0.09 rows=0) (actual time=0.004..0.004 rows=1 loops=2974)
|     -> Index lookup on f3 using foodName_idx (foodName=f1.foodName) (cost=0.09 rows=1) (actual time=0.003..0.003 rows=1 loops=2974)
|     -> Select #4 (subquery in condition; run only once)
|       -> Aggregate: avg(f4.proteinCalories) (cost=1476.85 rows=7263) (actual time=2.797..2.797 rows=1 loops=1)
|         -> Table scan on f4 (cost=750.55 rows=7263) (actual time=0.014..1.425 rows=7083 loops=1)
+-----+
```

If we add the carbsCal_idx the cost and rows increases

```
1 row in set (0.02 sec)
```

```
mysql> DROP INDEX foodName_idx on Food;
Query OK, 0 rows affected (0.03 sec)
Records: 0  Duplicates: 0  Warnings: 0
```

When we drop the foodname_idx, so we only have the carbsCal_idx, the cost and rows went back high.

Also if we replace carbsCalories by proteinCalories or fibreCalories it remains unchanged since the values for all three are same number of rows and the only difference is in the value of the integers so it does not account for any change in cost or number of rows used.

Thus, we know that the **drop of cost and row is driven by the foodName_idx index. So we are going with foodName_idx for this specific query for now.**

Index #2:

Without Indexing:

```
-----+
| -> Nested loop inner join (cost=586515.26 rows=5860069) (actual time=7.252..11.768 rows=1890 loops=1)
|   -> Filter: (f1.carbsCalories > (select #2)) (cost=266.33 rows=2421) (actual time=2.179..4.676 rows=2974 loops=1)
|     -> Table scan on f1 (cost=266.33 rows=7263) (actual time=0.043..1.697 rows=7083 loops=1)
|     -> Select #2 (subquery in condition; run only once)
|       -> Aggregate: avg(f2.carbsCalories) (cost=1476.85 rows=7263) (actual time=2.119..2.119 rows=1 loops=1)
|         -> Table scan on f2 (cost=750.55 rows=7263) (actual time=0.013..1.346 rows=7083 loops=1)
|   -> Single-row index lookup on <subquery3> using <auto_distinct_key> (foodName=f1.foodName) (actual time=0.000..0.000 rows=1 loops=2974)
|     -> Materialize with deduplication (cost=508.40..508.40 rows=2421) (actual time=6.478..6.657 rows=2270 loops=1)
|       -> Filter: (f3.proteinCalories > (select #4)) (cost=266.33 rows=2421) (actual time=1.935..4.226 rows=2270 loops=1)
|         -> Table scan on f3 (cost=266.33 rows=7263) (actual time=0.017..1.449 rows=7083 loops=1)
|         -> Select #4 (subquery in condition; run only once)
|           -> Aggregate: avg(f4.proteinCalories) (cost=1476.85 rows=7263) (actual time=1.907..1.907 rows=1 loops=1)
|             -> Table scan on f4 (cost=750.55 rows=7263) (actual time=0.015..1.232 rows=7083 loops=1)
|
+-----+
|
+-----+
1 row in set (0.02 sec)
```

Here, we can see that the cost of the nested loop inner join is 586515.26 and it has to get through 5860069 rows.

After Indexing:

```
mysql> CREATE INDEX two_idx on Food(foodName, carbsCalories);
Query OK, 0 rows affected (0.14 sec)
Records: 0  Duplicates: 0  Warnings: 0
```

After creating an index for the food name and carbsCalories (two_idx) the cost significantly drops to 1135.37 and it only has to go through 828 rows.

```
-> Nested loop semijoin (cost=1135.37 rows=828) (actual time=5.799..21.218 rows=1890 loops=1)
-> Filter: (f1.carbsCalories > (select #2)) (cost=266.33 rows=2421) (actual time=2.763..6.598 rows=2974 loops=1)
-> Table scan on f1 (cost=266.33 rows=7263) (actual time=0.043..2.200 rows=7083 loops=1)
-> Select #2 (subquery in condition; run only once)
-> Aggregate: avg(f2.carbsCalories) (cost=1476.85 rows=7263) (actual time=2.695..2.695 rows=1 loops=1)
-> Index scan on f2 using foodName_idx (cost=750.55 rows=7263) (actual time=0.015..1.311 rows=7083 loops=1)
-> Filter: (f3.proteinCalories > (select #4)) (cost=0.09 rows=0) (actual time=0.005..0.005 rows=1 loops=2974)
-> Index lookup on f3 using foodName_idx (foodName=f1.foodName) (cost=0.09 rows=1) (actual time=0.003..0.003 rows=1 loops=2974)
-> Select #4 (subquery in condition; run only once)
-> Aggregate: avg(f4.proteinCalories) (cost=1476.85 rows=7263) (actual time=3.000..3.000 rows=1 loops=1)
-> Table scan on f4 (cost=750.55 rows=7263) (actual time=0.015..1.610 rows=7083 loops=1)
```

Our second option can be to use `two_index` as our index since it significantly reduces our cost and number of rows. However, since it doesn't make any difference from using just `foodName_idx`, we would go ahead and choose it as the index.

Index #3:

Without Indexing:

[illegible]

Here, we can see that the cost of the nested loop inner join is 586515.26 and it has to get through 5860069 rows without indexing.

After Indexing:

```
mysql> CREATE INDEX three_idx on Food(foodName, carbsCalories, proteinCalories);
Query OK, 0 rows affected (0.13 sec)
Records: 0 Duplicates: 0 Warnings: 0
```

```
-> Remove duplicate f1 rows using temporary table (weedout) (cost=1135.37 rows=828) (actual time=5.619..19.752 rows=1890 loops=1)
-> Nested loop inner join (cost=1135.37 rows=828) (actual time=5.610..18.721 rows=1890 loops=1)
-> Filter: (f3.proteinCalories > (select #4)) (cost=266.33 rows=2421) (actual time=2.857..6.275 rows=2270 loops=1)
-> Index scan on f3 using foodName_idx (cost=266.33 rows=7263) (actual time=0.046..1.764 rows=7083 loops=1)
-> Select #4 (subquery in condition; run only once)
-> Aggregate: avg(f4.proteinCalories) (cost=1476.85 rows=7263) (actual time=2.800..2.800 rows=1 loops=1)
-> Index scan on f4 using foodName_idx (cost=750.55 rows=7263) (actual time=0.013..1.392 rows=7083 loops=1)
-> Filter: (f1.carbsCalories > (select #2)) (cost=0.26 rows=0) (actual time=0.005..0.005 rows=1 loops=2270)
-> Index lookup on f1 using foodName_idx (foodName=f3.foodName) (cost=0.26 rows=1) (actual time=0.003..0.004 rows=1 loops=2270)
-> Select #2 (subquery in condition; run only once)
-> Aggregate: avg(f2.carbsCalories) (cost=1476.85 rows=7263) (actual time=2.721..2.722 rows=1 loops=1)
-> Index scan on f2 using foodName_idx (cost=750.55 rows=7263) (actual time=0.014..1.340 rows=7083 loops=1)
```

When we index the foodName, carbsCalories, and proteinCalories, we find that it reduces the cost of the nested loop inner join significantly as compared to no indexing at all. However, the nested loop inner join cost remains the same that we get when we added index# 1 and index#2 which is 1135.37 and the rows 828. Moreover, it also adds on another cost of removing the duplicate rows using a temporary table which makes it inefficient as compared to what we get from index #1 and index#2.

Conclusion

Thus, we know that the drop is driven by indexing foodName. So we are going with indexing foodName for this specific query for now.

Advanced Query 2:

```
EXPLAIN ANALYZE
SELECT firstName, lastName, MAX(BMI) as maxBMI, MIN(BMI) as minBMI, AVG(BMI) as avgBMI
FROM health_record JOIN user ON (healthUserID = userID)
GROUP BY healthUserID
HAVING healthUserID = 1;
```

Index #1

Before indexing:

Before indexing our analysis for the table looks like this where the actual time for filter is 5.178..5.573 and the time for aggregate is 5.176..5.426

Creating an index on healthUserId as healthUserId_idx:

```
mysql> CREATE INDEX healthUserID_idx on health_record(healthUserID);
Query OK, 0 rows affected, 1 warning (0.06 sec)
Records: 0  Duplicates: 0  Warnings: 1
```

After indexing:

```
-> Filter: (health_record.healthUserID = 1) (actual time=4.951..5.334 rows=1 loops=1)
-> Table scan on <temporary> (actual time=0.001..0.161 rows=1000 loops=1)
-> Aggregate using temporary table (actual time=4.950..5.206 rows=1000 loops=1)
-> Nested loop inner join (cost=469.35 rows=1051) (actual time=0.076..3.308 rows=1051 loops=1)
-> Table scan on user (cost=101.50 rows=1000) (actual time=0.052..0.352 rows=1000 loops=1)
-> Index lookup on health_record using healthUserID (healthUserID='user'.userID) (cost=0.26 rows=1) (actual time=0.002..0.003
rows=1 loops=1000)
```

After indexing our actual time for filter becomes 4.951..5.344 and the actual time for aggregate becomes 4.950..5.206

We aren't not selecting h

Index #2:

Before indexing:

```
| -> Filter: (health_record.healthUserID = 1) (actual time=5.178..5.573 rows=1 loops=1)
-> Table scan on <temporary> (actual time=0.001..0.154 rows=1000 loops=1)
-> Aggregate using temporary table (actual time=5.176..5.426 rows=1000 loops=1)
-> Nested loop inner join (cost=469.35 rows=1051) (actual time=0.069..3.397 rows=1051 loops=1)
-> Table scan on user (cost=101.50 rows=1000) (actual time=0.046..0.351 rows=1000 loops=1)
-> Index lookup on health_record using healthUserID (healthUserID='user'.userID) (cost=0.26 rows=1) (actual time=0.002..0.003
rows=1 loops=1000)
```

Before indexing our analysis for the table looks like this where the actual time for filter is 5.178..5.573 and the time for aggregate is 5.176..5.426

Creating index on first name as firstName_idx:

```
mysql> CREATE INDEX firstName_idx on user(firstName);
Query OK, 0 rows affected (0.05 sec)
Records: 0  Duplicates: 0  Warnings: 0
```

After indexing:


```
| -> Filter: (health_record.healthUserID = 1) (actual time=5.835..6.227 rows=1 loops=1)
    -> Table scan on <temporary> (actual time=0.001..0.153 rows=1000 loops=1)
        -> Aggregate using temporary table (actual time=5.831..6.080 rows=1000 loops=1)
            -> Nested loop inner join (cost=469.35 rows=1051) (actual time=0.047..3.773 rows=1051 loops=1)
                -> Table scan on user (cost=101.50 rows=1000) (actual time=0.032..0.407 rows=1000 loops=1)
                -> Index lookup on health_record using healthUserID (healthUserID='user'.userID) (cost=0.26 rows=1) (actual time=0.002..0.003 rows=1 loops=1000)
```

After indexing our actual time for filter becomes 5.835..6.227 and the actual time for aggregate becomes 5.831..6.080 which is greater than the original by a small amount

Since this index increases our time for filter and aggregate we won't be using firstName as our index for this query and firstName is a primary key so creating an index for that won't be of any use.

Index#3:

Before indexing:

```
| -> Filter: (health_record.healthUserID = 1) (actual time=5.178..5.573 rows=1 loops=1)
    -> Table scan on <temporary> (actual time=0.001..0.154 rows=1000 loops=1)
        -> Aggregate using temporary table (actual time=5.176..5.426 rows=1000 loops=1)
            -> Nested loop inner join (cost=469.35 rows=1051) (actual time=0.069..3.397 rows=1051 loops=1)
                -> Table scan on user (cost=101.50 rows=1000) (actual time=0.046..0.351 rows=1000 loops=1)
                -> Index lookup on health_record using healthUserID (healthUserID='user'.userID) (cost=0.26 rows=1) (actual time=0.002..0.003 rows=1 loops=1000)
```

Before indexing our analysis for the table looks like this where the actual time for filter is 5.178..5.573 and the time for aggregate is 5.176..5.426

Creating index on firstName and lastName as name_idx:

```
mysql> CREATE INDEX name_idx on user(firstName,lastName);
Query OK, 0 rows affected (0.05 sec)
Records: 0  Duplicates: 0  Warnings: 0
```

After indexing:

```
| -> Filter: (health_record.healthUserID = 1) (actual time=5.024..5.367 rows=1 loops=1)
    -> Table scan on <temporary> (actual time=0.001..0.152 rows=1000 loops=1)
        -> Aggregate using temporary table (actual time=4.935..5.182 rows=1000 loops=1)
            -> Nested loop inner join (cost=469.35 rows=1051) (actual time=0.041..3.318 rows=1051 loops=1)
                -> Index scan on user using name_idx (cost=101.50 rows=1000) (actual time=0.029..0.335 rows=1000 loops=1)
                -> Index lookup on health_record using healthUserID (healthUserID='user'.userID) (cost=0.26 rows=1) (actual time=0.002..0.003 rows=1 loops=1000)
```

After indexing our actual time for filter becomes 5.024..5.367 and the actual time for aggregate becomes 4.935..5.182 which provides us with a small change but which is

better than our first index. However, creating index on firstName and lastName will be of no use as they are primary keys.

Index#4:

Before indexing

```
mysql> EXPLAIN ANALYZE
-> SELECT firstName, lastName, MAX(BMI) as maxBMI, MIN(BMI) as minBMI, AVG(BMI) as avgBMI
-> FROM health_record JOIN user ON (healthUserID = userID)
-> GROUP BY healthUserID
-> HAVING healthUserID = 1;
+-----+
| EXPLAIN |
+-----+
| -> Filter: (health_record.healthUserID = 1) (actual time=203.624..203.836 rows=1 loops=1)
| -> Table scan on <temporary> (actual time=0.001..0.101 rows=1000 loops=1)
| -> Aggregate using temporary table (actual time=203.620..203.778 rows=1000 loops=1)
| -> Nested loop inner join (cost=1206.00 rows=1000) (actual time=148.464..202.468 rows=1051 loops=1)
| -> Table scan on user (cost=106.00 rows=1000) (actual time=46.753..93.146 rows=1000 loops=1)
| -> Index lookup on health_record using healthUserID (healthUserID='user'.userID) (cost=1.00 rows=1) (actual time=0.108..0.109 rows=1 loops=1000)
+-----+
1 row in set (0.20 sec)
```

Create index

```
mysql> CREATE INDEX BMI_idx on health_record(BMI);
Query OK, 0 rows affected (0.10 sec)
Records: 0 Duplicates: 0 Warnings: 0
```

After indexing

```
mysql> EXPLAIN ANALYZE SELECT firstName, lastName, MAX(BMI) as maxBMI, MIN(BMI) as minBMI, AVG(BMI) as avgBMI FROM health_record JOIN user ON (healthUserID = userID) GROUP BY healthUserID HAVING healthUserID = 1;
+-----+
| EXPLAIN |
+-----+
| -> Filter: (health_record.healthUserID = 1) (actual time=3.908..4.113 rows=1 loops=1)
| -> Table scan on <temporary> (actual time=0.001..0.085 rows=1000 loops=1)
| -> Aggregate using temporary table (actual time=3.906..4.051 rows=1000 loops=1)
| -> Nested loop inner join (cost=451.50 rows=1000) (actual time=0.094..2.944 rows=1051 loops=1)
| -> Table scan on user (cost=101.50 rows=1000) (actual time=0.049..0.349 rows=1000 loops=1)
| -> Index lookup on health_record using healthUserID (healthUserID='user'.userID) (cost=0.25 rows=1) (actual time=0.002..0.002 rows=1 loops=1000)
+-----+
1 row in set (0.01 sec)
```

After indexing BMI, the actual time for the filter becomes 3.908 and the actual time for aggregate is 3.906. Thus, we can see that out of all the indexes and without indexing the best result is provided by indexing BMI. It improves the performance of the data

matching by reducing the time taken to match the query values for BMI. Hence, we would be using this index for the above listed query.

Conclusion

Hence, we can see that the most effective index was after indexing BMI which would be the best choice for us to use for the above 2nd advanced query as it helped reduce the time significantly and is the most efficient one out of all the other.

Latest SQL Code:

https://s3-us-west-2.amazonaws.com/secure.notion-static.com/8338eae-a39d-4d82-908a-2fd20b585/database_latest.sql